

Lecture 8

SQL part I

Subject Lecturer: Kevin K.F. YUEN, PhD.

Acknowledgement: Slides were offered from Prof. Ken Yiu.
Some parts might be revised and indicated.

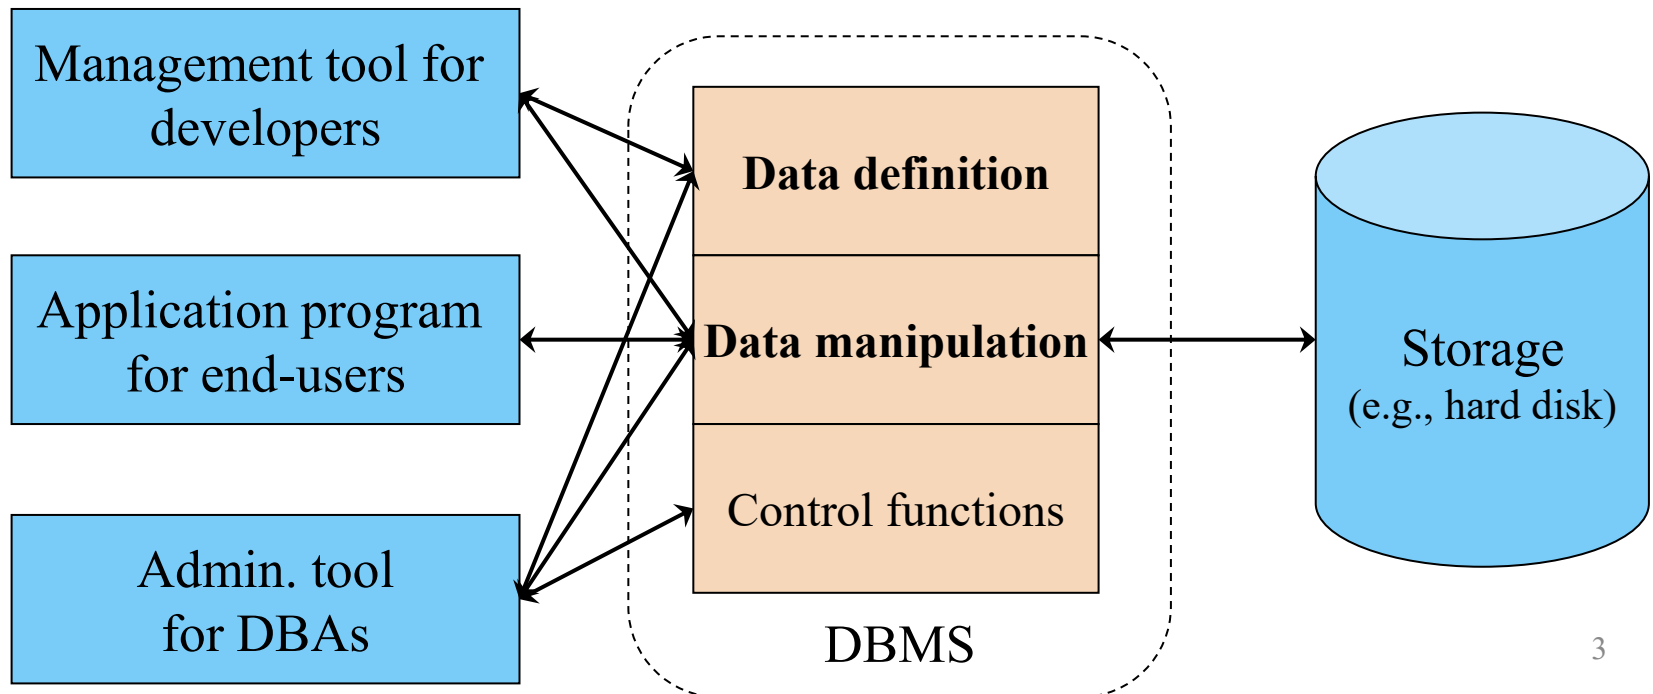
Outline



- ◆ Background
- ◆ How to define a table?
- ◆ Basic SQL: select, from, where
- ◆ Order by, aggregation, group by
- ◆ Set operations
- ◆ How to insert, delete, update data?

How to use DBMS?

- ◇ *Data Definition Language (DDL)*:
 - ◆ define the database schema, i.e., the record type
- ◇ *Data Manipulation Language (DML)*:
 - ◆ access / update the database content
- ◇ SQL consists of both DDL and DML



History of SQL

- ◆ Early days (1970s)
 - ◆ Known as SEQUEL (Structured English Query Language)
 - ◆ Developed for IBM's system R
- ◆ Now called SQL (Structured Query Language)
- ◆ Standardization by ISO
 - ◆ SQL-86
 - ◆ SQL-89
 - ◆ + integrity constraints, 100 pages
 - ◆ SQL-92
 - ◆ + new DDL,DML features, 500 pages
 - ◆ SQL:1999, 2003, 2006, 2008, 2011, 2016
 - ◆ + many features, many pages

DBMS Implementations

- ◆ DBMS implementations
 - ◆ Microsoft SQL server, Oracle, IBM DB2,
- ◆ Many implementations fully support “intermediate SQL” --- half of new features in SQL-92
- ◆ Some DBMS vendors have proprietary extensions to the SQL standard
 - ◆ E.g., a SQL statement used in DBMS “A” may not be directly supported in DBMS “B”

SQL vs. relational algebra

- ◆ In SQL, **duplicates** are allowed in a table
 - ◆ In relational algebra, duplicates are not allowed
- ◆ Some functions are supported in SQL but not in relational algebra
 - ◆ Data definition language (e.g., create table, drop table)
 - ◆ Nested queries
 - ◆ Sorting a table
 - ◆ Keywords for data types
 - ◆ Conditional expressions (CASE)
 - ◆ String functions

Outline

- ◆ Background



- ◆ How to define a table?
- ◆ Basic SQL: select, from, where
- ◆ Order by, aggregation, group by
- ◆ Set operations
- ◆ How to insert, delete, update data?

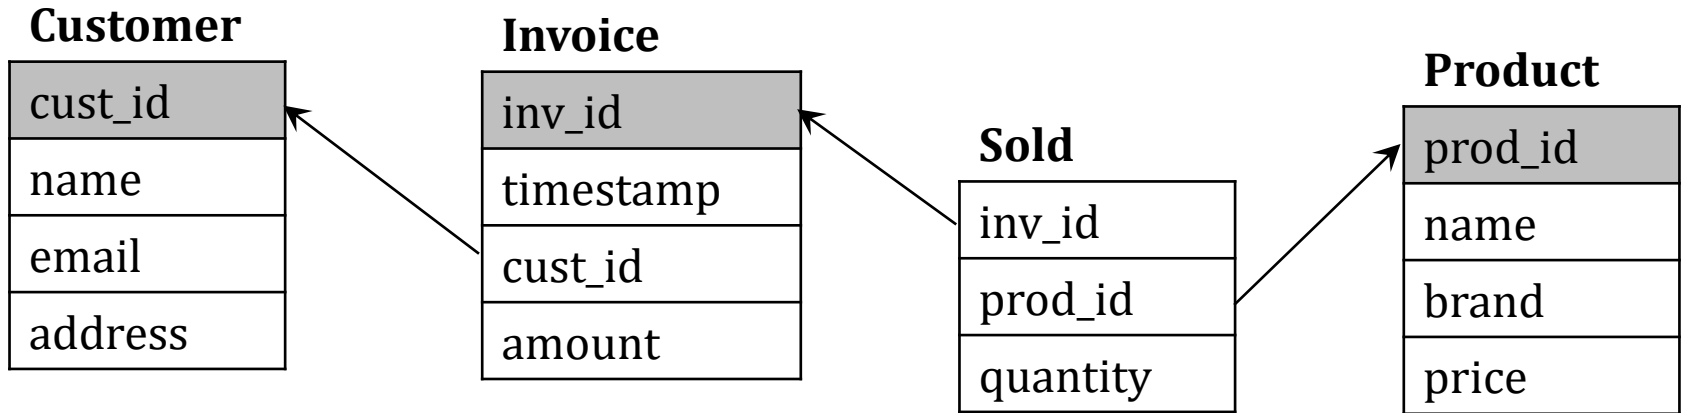
How to create tables?

- ◆ When we create a table, we provide:
 - ◆ The table name
 - ◆ The name and type of each attribute name
 - ◆ A name may contain '_' but not '-'
 - ◆ Integrity constraints
 - ◆ E.g., primary key, foreign key, NOT NULL, unique key, CHECK
- ◆ Basic types
 - ◆ **char**(n): a string with fixed length n
 - ◆ **int**: an integer
 - ◆ **numeric**(p, d): a fixed point number with total p digits and d digits after the decimal point
 - ◆ **real**: floating-point number

Customer

cust_id
name
email
address

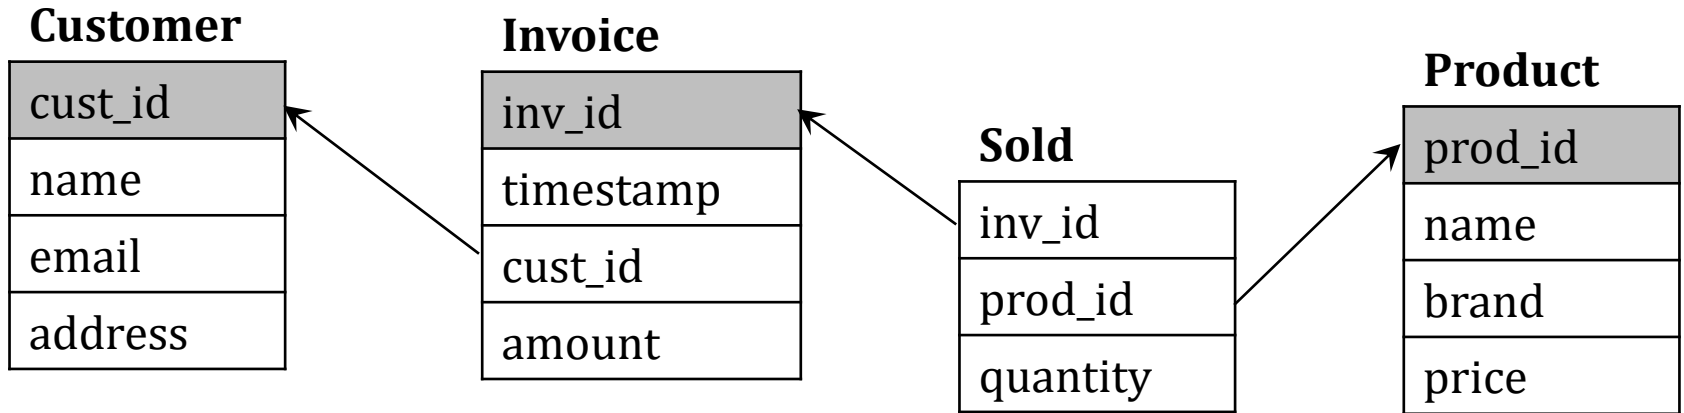
How to create tables?



◆ Create the Invoice table by:

```
create table Invoice
(inv_id int,
timestamp int,
cust_id int,
amount numeric(10,2),
primary key (inv_id))
```

How to create tables?



- ◆ Create the Sold table by:


```
create table sold
(inv_id int,
prod_id int,
quantity int,
primary key (inv_id,prod_id))
```

Other SQL statements for data definition

Purpose	Example statement
Show tables in a database	<code>show tables</code>
Describe the schema of a table	<code>describe Invoice</code>
Remove a table	<code>drop table Invoice</code>
Add a column in a table	<code>alter table Invoice add staff_id int</code>
Drop a column in a table	<code>alter table Invoice drop staff_id</code>

```
+-----+-----+-----+-----+-----+
| Field      | Type          | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+
| inv_id     | int(11)       | NO   | PRI | 0        |       |
| timestamp  | int(11)       | YES  |     | NULL     |       |
| cust_id    | int(11)       | YES  |     | NULL     |       |
| amount     | decimal(10,2) | YES  |     | NULL     |       |
+-----+-----+-----+-----+-----+
4 rows in set (0.01 sec)
```

Outline

- ◆ Background
- ◆ How to define a table?
-  ◆ Basic SQL: select, from, where
- ◆ Order by, aggregation, group by
- ◆ Set operations
- ◆ How to insert, delete, update data?

SQL and Relational Algebra

- ◆ A simple SQL query has the form:

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

- ◆ A_i represents an attribute
- ◆ r_i represents a relation
- ◆ P is a predicate

- ◆ Equivalent to the relational algebra expression

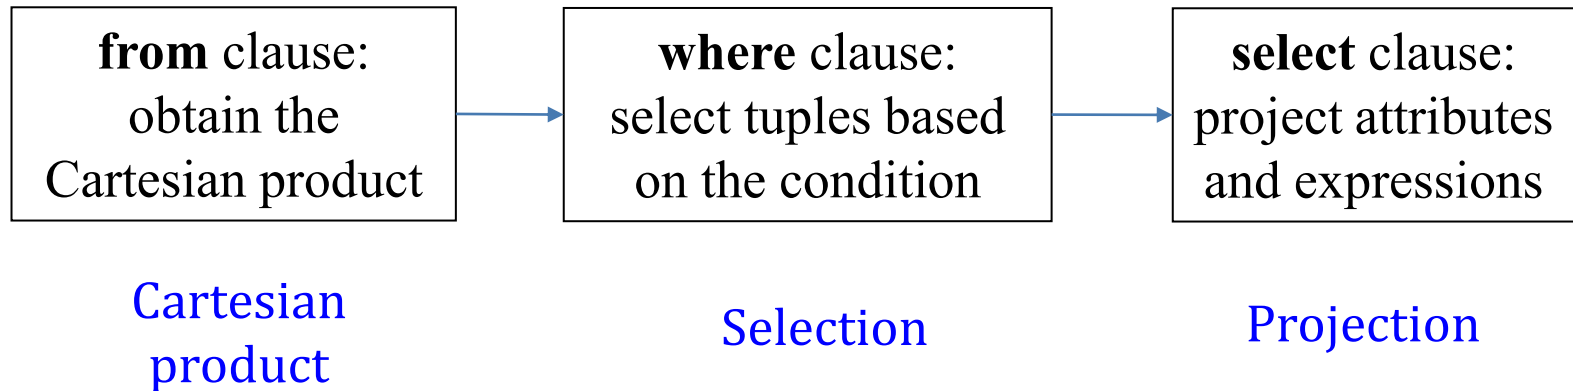
Projection Selection Cartesian product

$$\Pi_{A_1, A_2, \dots, A_n} (\sigma_P (r_1 \times r_2 \times \dots \times r_m))$$

- ◆ The result of an SQL query is a relation
 - ◆ Note: the **select** keyword in SQL means projection (but not selection) in relational algebra!

Flow of Clauses

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P



How to express the following queries in SQL?

- ◆ (Q1) Find **products** that have **price** > 8.0
- ◆ (Q2) Find the **brands** of **products**
- ◆ (Q3) Find the **customer name** and **inv_id** (invoice id) of each **invoice**
- ◆ (Q4) Find the highest **price** in the table *Product*
- ◆ (Q5) Find the total **amount** in the table *Invoice*

SQL: where clause

- ◆ (Q1) Find products that have price > 8.0

select *

* means all attributes

from Product

where price > 8.0

Relation **Product**:

prod_id	name	brand	price
1	Coca Cola	CO	7.8
2	Pepsi	PE	8.9
3	7 Up	DP	6.5
4	Sprite	CO	8.3



prod_id	name	brand	price
2	Pepsi	PE	8.9
4	Sprite	CO	8.3

SQL: where clause

- ◆ In the condition, we may use
 - ◆ Comparisons: =, !=, <, >, <=, >=
 - ◆ Connectives: and, or, not, between...and...

- ◆ Examples:

```
select *  
from Product  
where price>=8.0 and price<=9.0
```

```
select *  
from Product  
where price between 8.0 and 9.0
```

SQL: select clause

- ❖ (Q2) Find the brands of products
 - ❖ By default, SQL allows duplicates
 - ❖ To remove duplicates, use the keyword **distinct**

select brand
from Product

brand
CO
PE
DP
CO

select distinct brand
from Product

brand
CO
PE
DP

prod_id	name	brand	price
1	Coca Cola	CO	7.8
2	Pepsi	PE	8.9
3	7 Up	DP	6.5
4	Sprite	CO	8.3

SQL: from clause, rename

- ❖ (Q3) “find the customer name and inv_id of each invoice”
 - ❖ Place both tables Customer and Invoice in the from clause
 - ❖ *Optional*: use the keyword **as** to rename a table and make the SQL statement easier to read

Remove 'as' in oracle


```
select I.inv_id, C.name
from Customer as C, Invoice as I
where C.cust_id=I.cust_id
```

cust_id	name	email	address
1	James	james@yahoo.com	AB
2	Mary	mary@gmail.com	CD
3	Peter	peter@yahoo.com	EF
4	Peter	peter@gmail.com	null

inv_id	timestamp	cust_id	amount
1	101	3	8.9
2	102	2	7.8

inv_id	name
1	Peter
2	Mary

Outline

- ◆ Background
- ◆ How to define a table?
- ◆ Basic SQL: select, from, where
-  ◆ Order by, aggregation, group by
- ◆ Set operations
- ◆ How to insert, delete, update data?

SQL: order by

- ◆ The order by clause allows us to sort the results
 - ◆ **asc**: ascending order; **desc**: descending order

```
select *  
from Product  
order by price asc
```



prod_id	name	brand	price
3	7 Up	DP	6.5
1	Coca Cola	CO	7.8
4	Sprite	CO	8.3
2	Pepsi	PE	8.9

```
select *  
from Product  
order by price desc
```



prod_id	name	brand	price
2	Pepsi	PE	8.9
4	Sprite	CO	8.3
1	Coca Cola	CO	7.8
3	7 Up	DP	6.5

SQL: aggregation

- ◆ Aggregate functions
 - ◆ sum, count, avg, min, max
- ◆ (Q4) Find the highest **price** in the table *Product*
- ◆ (Q5) Find the total **amount** in the table *Invoice*

select max(price)
from Product

select sum(amount)
from Invoice

inv_id	timestamp	cust_id	amount
1	101	3	8.9
2	102	2	7.8
3	103	2	6.5
4	104	3	8.3



sum(amount)
31.5

SQL: group by

- ◆ E.g., find the total amount spent by each customer
 - ◆ First partition the table into **groups** (by cust_id), then use aggregate function on each **group**

```
select cust_id, sum(amount)
from Invoice
group by cust_id
```

inv_id	timestamp	cust_id	amount
1	101	3	8.9
2	102	2	7.8
3	103	2	6.5
4	104	3	8.3

group by



inv_id	timestamp	cust_id	amount
1	101	3	8.9
4	104	3	8.3
2	102	2	7.8
3	103	2	6.5

aggregation



cust_id	sum(amount)
3	17.2
2	14.3

SQL: group by + having

- ◆ Keyword **having**: specify filter condition for groups
- ◆ E.g., find the total amount spent by each customer; display those with total amount greater than 15.0

```
select cust_id, sum(amount)
from Invoice
group by cust_id
having sum(amount)>15.0
```

inv_id	timestamp	cust_id	amount
1	101	3	8.9
2	102	2	7.8
3	103	2	6.5
4	104	3	8.3

group by



inv_id	timestamp	cust_id	amount
1	101	3	8.9
4	104	3	8.3
2	102	2	7.8
3	103	2	6.5

aggregation



cust_id	sum(amount)
3	17.2

having



cust_id	sum(amount)
3	17.2
2	14.3

Outline

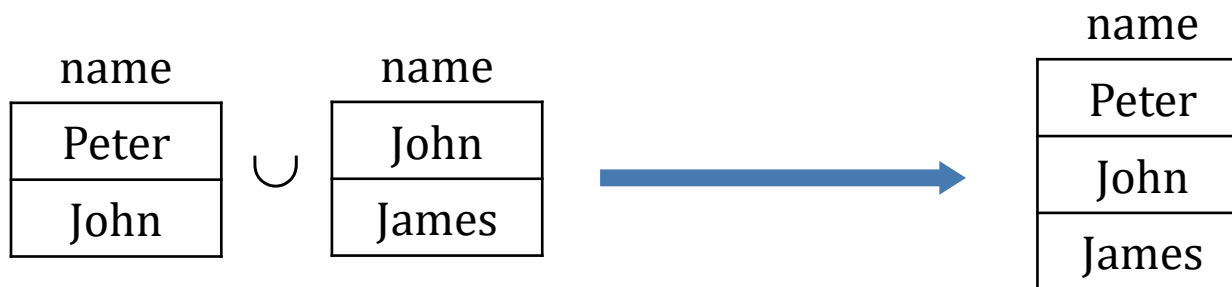
- ◆ Background
- ◆ How to define a table?
- ◆ Basic SQL: select, from, where
- ◆ Order by, aggregation, group by
- ◆ Set operations
- ◆ How to insert, delete, update data?



SQL: set operations

- ◆ Set operations: **union**, **intersect**, **except**
 - ◆ By default, these operations **remove duplicates**

```
(select name from A)  
union  
(select name from B)
```



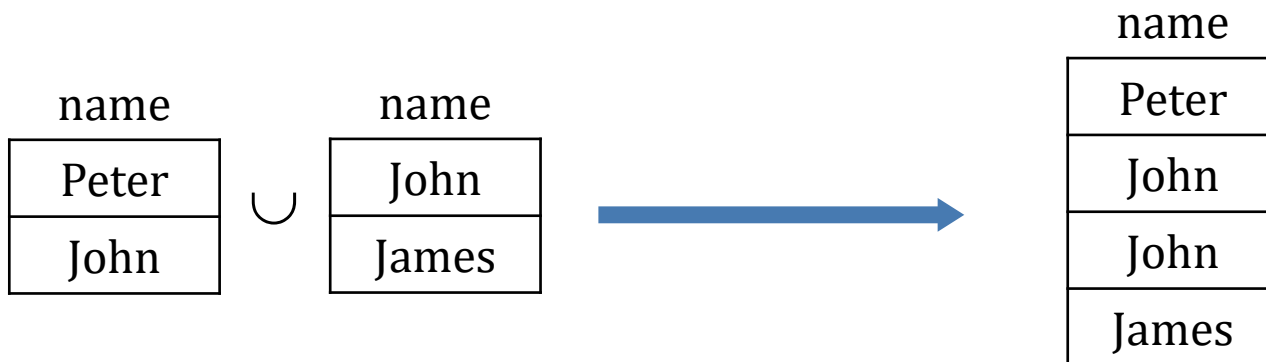
SQL: set operations

Minus in Oracle

MS SQL

- ◆ Set operations: **union**, **intersect**, **except**
 - ◆ To allow duplicates, use the keyword **all**


```
(select name from A)
union all
(select name from B)
```



Oracle MINUS is an operator; it's equivalent to EXCEPT in SQL Server.

<https://stackoverflow.com/questions/5557991/minus-vs-except-difference-in-oracle-sql-server>

Outline

- ◆ Background
- ◆ How to define a table?
- ◆ Basic SQL: select, from, where
- ◆ Order by, aggregation, group by
- ◆ Set operations
-  ◆ How to insert, delete, update data?

Insertion

- ◆ **Insertion:** insert tuple(s) into a relation

- ◆ Example 1: insert a tuple

```
insert into Product  
values (5, 'Soda', 'AB', 7.5)
```

prod_id	name	brand	price
1	Coca Cola	CO	7.8
2	Pepsi	PE	8.9
4	Sprite	CO	8.3
5	Soda	AB	7.5

- ◆ Example 2: insert tuples obtained by select statement

```
insert into Product  
select ...  
from ...  
where ...
```

Deletion

- ◆ **Deletion:** delete tuple(s) from a relation
 - ◆ Example: remove the product that has id=3

delete from Product
where prod_id=3

prod_id	name	brand	price
1	Coca Cola	CO	7.8
2	Pepsi	PE	8.9
3	7-Up	DP	6.5
4	Sprite	CO	8.3

Update

- ◆ Update tuple(s) based on a condition
 - ◆ Example: double the price of each product of the brand 'AB'

```
update Product  
set price=price*2  
where brand='AB'
```

prod_id	name	brand	price
1	Coca Cola	CO	7.8
2	Pepsi	PE	8.9
4	Sprite	CO	8.3
5	Soda	AB	7.5 15

Summary

- ◆ After this lecture, you should be able to:
 - 1) Apply SQL to define tables
 - 2) Apply SQL to express simple queries
- ◆ Please read Chapter 3 in the book
“*Database System Concepts*”, 7th Edition
- ◆ Next lecture: advanced features of SQL