

Lecture 13

Query processing

Subject Lecturer: Kevin K.F. YUEN, PhD.

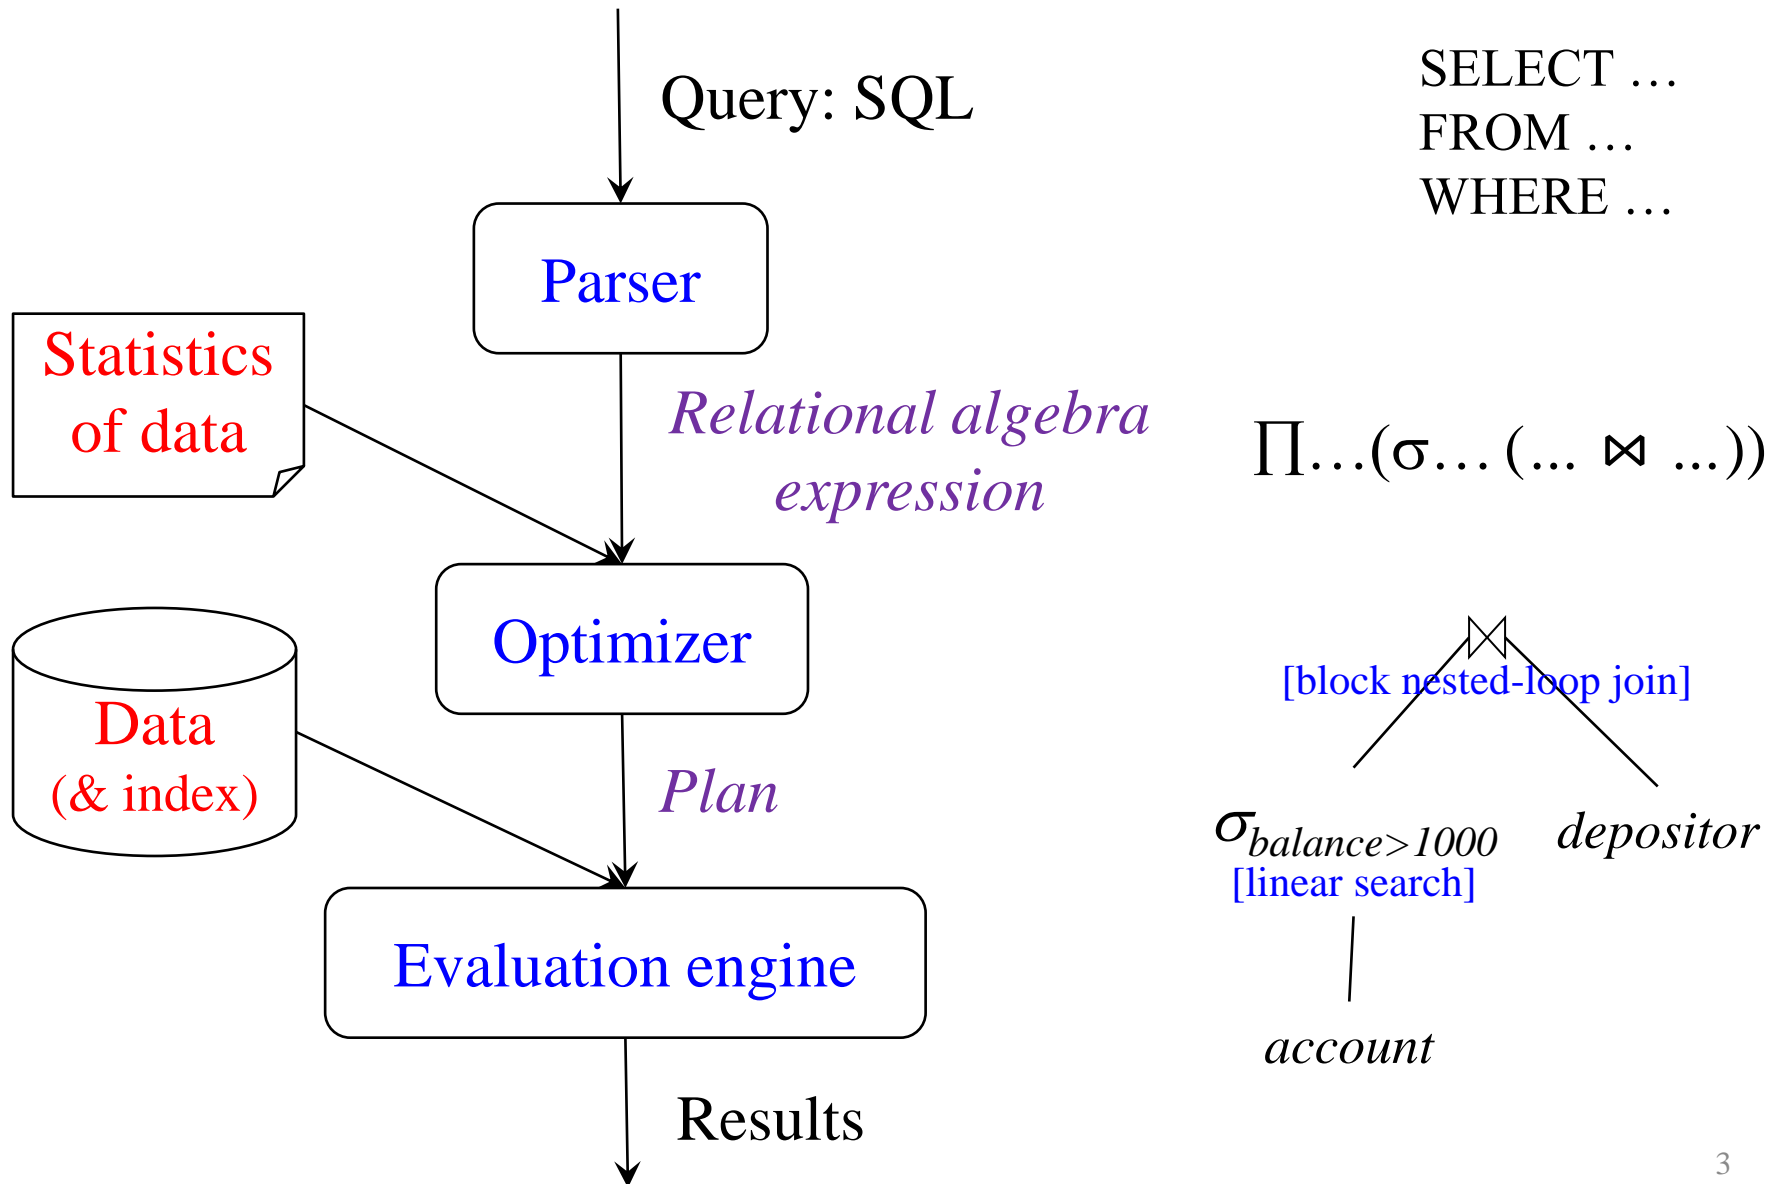
Acknowledgement: Slides were offered from Prof. Ken Yiu.
Some parts might be revised and indicated.

Outline



- ◆ Basic concepts for query processing
- ◆ How to process a selection?
- ◆ How to process a join?
- ◆ How to execute a plan?

How to process a query?

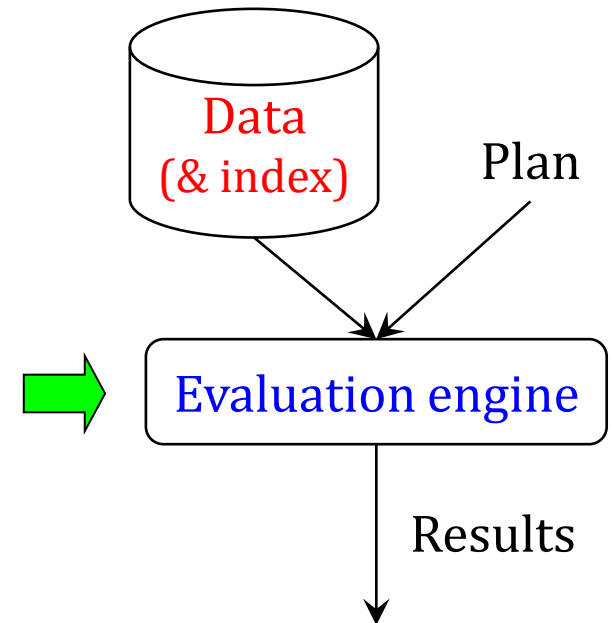


How to process a query?

- ◆ Optimizer
 - ◆ Find the evaluation plan with the lowest estimated cost
 - ◆ We skip this issue in this course
- ◆ Evaluation engine
 - ◆ Call an algorithm to evaluate a relational algebra operation
 - ◆ Combine individual operations to evaluate a complete plan

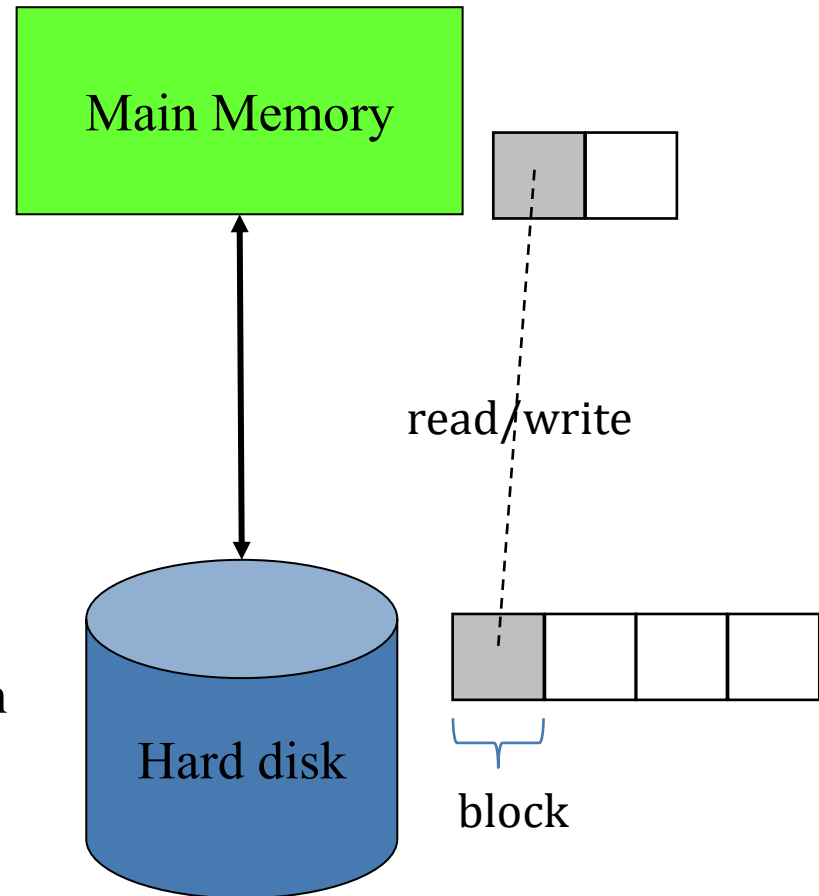
Evaluation engine: How to execute a plan?

- ◆ How to measure the cost of a plan?
- ◆ Some methods (physical operators) for executing
 - ◆ The selection operation
 - ◆ The join operation
- ◆ How to execute a plan?



Cost

- ◆ Cost = the number of disk block transfers
- ◆ Assumptions in RDBMS
 - ◆ Ignore CPU costs
 - ◆ Ignore the cost of writing the final output to disk
- ◆ Extra assumption in this lecture
 - ◆ Ignore the disk seek time, because the number of disk block transfers is much larger than the number of disk seeks



Outline

- ◆ Basic concepts for query processing



- ◆ How to process a selection?

- ◆ How to process a join?

- ◆ How to execute a plan?

Selection Operation

- ◆ *Example:* $\sigma_{\text{balance} < 2500}(\text{account})$
- ◆ Several different algorithms to implement selections
 - ◆ Usually choose the cheapest available one

Algorithm / physical operator	Cost (# disk blocks)
Linear search	b_r
Primary index, equality on key	$h_r + 1$
Primary index, equality on non-key	$h_r + b_{\text{results}}$
Secondary index, equality	$h_r + n_{\text{results}}$



b_r : size of r in blocks

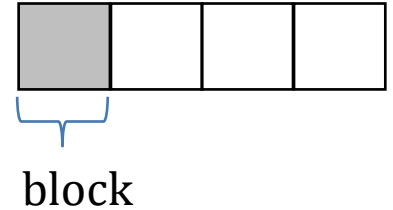
h_r : height of B⁺-tree on r

n_{results} : # of results

b_{results} : result size in blocks

For simplicity, we measure the cost as the number of disk block transfers.
 When using an index other than B⁺-tree, replace the term h_r by the index lookup cost.

Selection



◆ (1) *Linear search*:

Scan each file block and test all tuples

- ◆ Applicable to any type of comparison condition
- ◆ Cost = b_r blocks

(b_r : number of blocks occupied by relation r)

◆ (2) *Primary index on candidate key, equality*:

Retrieve a single tuple that satisfies the equality condition


- ◆ Cost = $h_r + 1$ blocks

- ◆ B⁺-tree height $h_r = \lceil \log_{\lceil f/2 \rceil} n_r \rceil$

- ◆ n_r : number of records in relation r

- ◆ f : max. number of children in a node

Selection



- ◆ (3) *Primary index on nonkey, equality*:

Retrieve multiple (consecutive) tuples that satisfies the equality condition

- ◆ $\text{Cost} = h_r + b_{match}$ blocks

- ◆ This requires estimating b_{match} : number of blocks containing matching tuples


- ◆ (4) *Secondary index on nonkey, equality* :

Retrieve multiple tuples that satisfies the equality condition

- ◆ $\text{Cost} = h_r + n_{match}$ blocks


- ◆ n_{match} denotes the number of matching records

Outline

- ◆ Basic concepts for query processing
- ◆ How to process a selection?
-  ◆ How to process a join?
- ◆ How to execute a plan?

Join Operation

- ◆ *Example:* **customer** ⋈ **depositor**
- ◆ Several different algorithms to implement joins
 - ◆ Usually choose the cheapest available one

Algorithm / physical operator	 Cost (# disk blocks)
Nested-loop join	$n_r * b_s + b_r$ [worst case]
Block nested-loop join	$\lceil b_r / (M-2) \rceil * b_s + b_r$
Indexed nested-loop join	$n_r * (h_s + 1) + b_r$
Merge-join	$ \begin{aligned} &b_r + b_s \\ &+ b_r (2 \lceil \log_{M-1}(b_r/M) \rceil + 1) \\ &+ b_s (2 \lceil \log_{M-1}(b_s/M) \rceil + 1) \end{aligned} $
Hash-join	$3(b_r + b_s)$ if no recursive partitioning required

For simplicity, we measure the cost as the number of disk block transfers.

Nested-Loop Join

for each tuple t_r in r do

for each tuple t_s in s do

if pair (t_r, t_s) satisfies the join condition

then add (t_r, t_s) to the result

◆ Applicable to any join condition, index not required

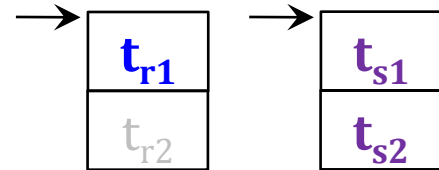
◆ Cost of *nested-loop join*:

$n_r * b_s + b_r$ blocks




◆ Assume the worst case: only one memory buffer block for each relation

Memory, 2 blocks



Disk



The sequence of blocks read from the disk



Index Nested-Loop Join

for each tuple t_r **in** r **do**

search the index on s to find tuples that match with t_r

for each matching tuple t_s **in** s **do**

if pair (t_r, t_s) satisfies the join condition

then add (t_r, t_s) to the result

- ◆ Requires an index, applicable to equality condition only
- ◆ Cost of *indexed nested-loop join*:
 $n_r * (h_s + 1) + b_r$ blocks

❖ Exercise: *customer* ⋈ *depositor*

- ❖ Number of tuples: $n_{customer}=10000$ $n_{depositor}=5000$
- ❖ Number of blocks: $b_{customer}=400$ $b_{depositor}=100$
- ❖ Suppose that *customer* has a primary B⁺-tree index on the join attribute *customer-name*, which contains 20 child pointers per index node.

❖ Cost of *nested-loop join*

If outer relation = *depositor*

cost : _____ * _____ + _____ = _____ blocks

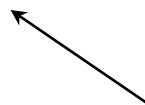
If outer relation = *customer*

cost : _____ * _____ + _____ = _____ blocks

❖ Cost of *indexed nested-loop join*

E.g., the tree height (for *customer*) is 4, so the cost is:

$$5000 * (\textcolor{red}{4} + 1) + 100 = 25100 \text{ blocks}$$



$$\lceil \log_{\lceil f/2 \rceil}(n_{customer}) \rceil$$
$$= \lceil \log_{\lceil 20/2 \rceil}(10000) \rceil = 4$$

* Try to find “log(10000)/log(10)”

How to compute this in calculator?

* Type “10000” “log” “/” “10” “log”

Block Nested-Loop Join

- Variant of nested-loop join
- Pair every block of the inner relation with every block of the outer relation

In this method, the relation s is scanned times

for each block B_r of r do



for each block B_s of s do

for each tuple t_r in B_r do

for each tuple t_s in B_s do

if pair (t_r, t_s) satisfies the join condition

then add (t_r, t_s) to the result



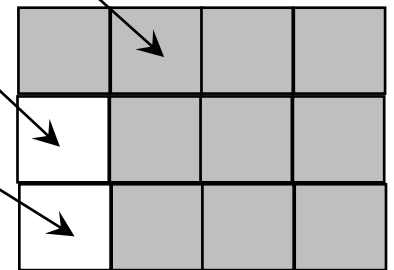
The sequence of blocks read from the disk



Block Nested-Loop Join (Cont.)

- ◆ More memory blocks may be used to reduce the cost of block nested-loop join
- ◆ If the memory has M blocks, use them as follows:
 - ◆ $M - 2$ memory blocks to buffer the outer relation
 - ◆ 1 block to buffer the inner relation
 - ◆ 1 block to buffer the output

- ◆ $\text{Cost} = \lceil b_r / (M - 2) \rceil * b_s + b_r$ blocks



memory

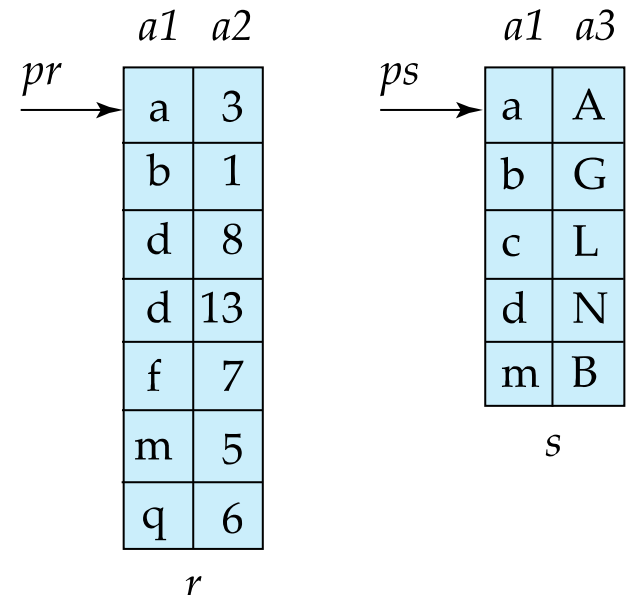
Merge-Join

1. Sort both relations on their join attribute (*to discuss soon*)
2. Merge the sorted relations to join them
 - a. This step is like the merge stage of the sort-merge algorithm.
 - b. The difference is to handle duplicate values in join attribute
— every pair with same value on join attribute must be matched

- ◆ Applicable to equi-joins and natural joins
- ◆ Each block is only read once
 - ◆ assuming all tuples for any given value of the join attributes fit in memory

- ◆ **Cost** of merge join =
 $b_r + b_s$ blocks

+ the sorting cost (if relations are unsorted)



Sorting: External Sort-Merge

Example: $M = 3$
memory blocks

Use it when the relation is larger than the main memory, i.e., $b_r > M$

External sort-merge algorithm

1. Create sorted runs
 - Read consecutive M blocks into memory, sort it, then write to a run
2. Merging until only 1 run remains

Cost: $b_r (2 \lceil \log_{M-1}(b_r/M) \rceil + 1)$

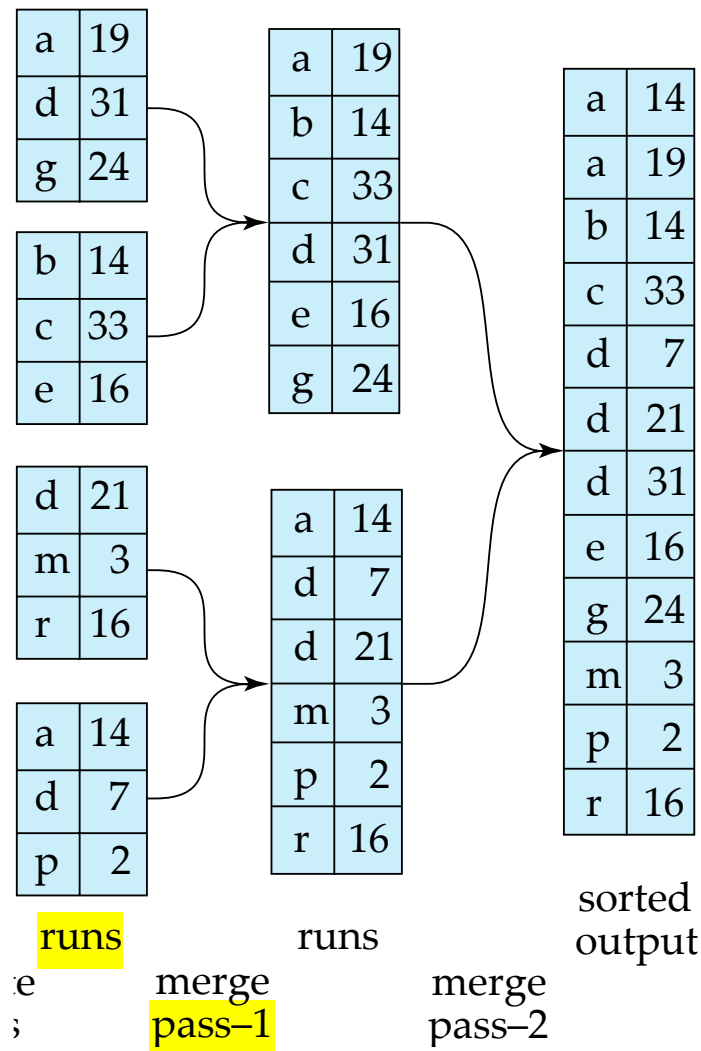
- Number of merge passes required: $\lceil \log_{M-1}(b_r/M) \rceil$
- Block transfers for initial run and in each pass is $2b_r$

initial relation	g	24	create runs	a	19
	a	19		d	31
	d	31		g	24
	c	33		b	14
	b	14		c	33
	e	16		e	16
	r	16		d	21
	d	21		m	3
	m	3		r	16
	p	2		a	14
	d	7		d	7
	a	14		p	2

Sorting: External Sort-Merge

- How to merge “sorted runs”?
- 2. Merging (for every $M-1$ runs)
 - Use $M-1$ memory blocks as input buffers, and 1 memory block as output buffer
 - Move the smallest tuple from its input buffer to the output buffer
 - An input block empty \rightarrow fill it with the next disk block from its input run
 - An output block full \rightarrow flush the block to its output run
- Repeat until only 1 run remains

Example: $M = 3$ memory blocks



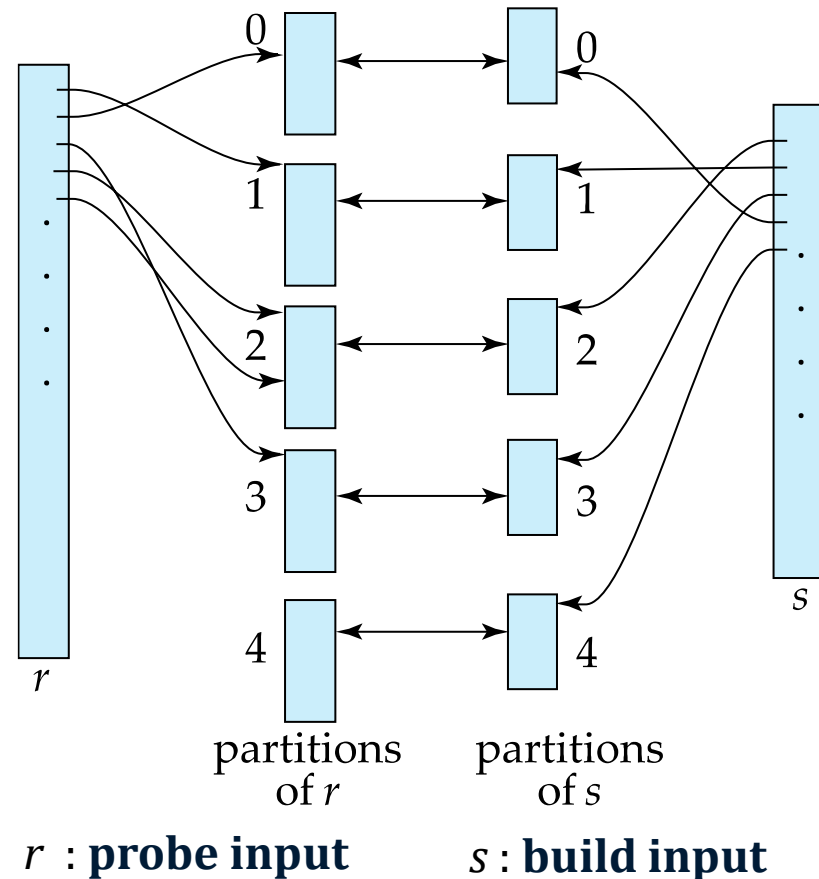
Hash-Join

Example: $n_h = 5$ partitions

Observation: if two records have the same key value (on the join attribute), then those two keys must have the same hash value

Idea of hash-join: use a hash function h to hash records into partitions

Hash-join is applicable to equi-joins and natural joins



Hash-Join

Example: $n_h = 5$ partitions

Call the smaller relation s as build input

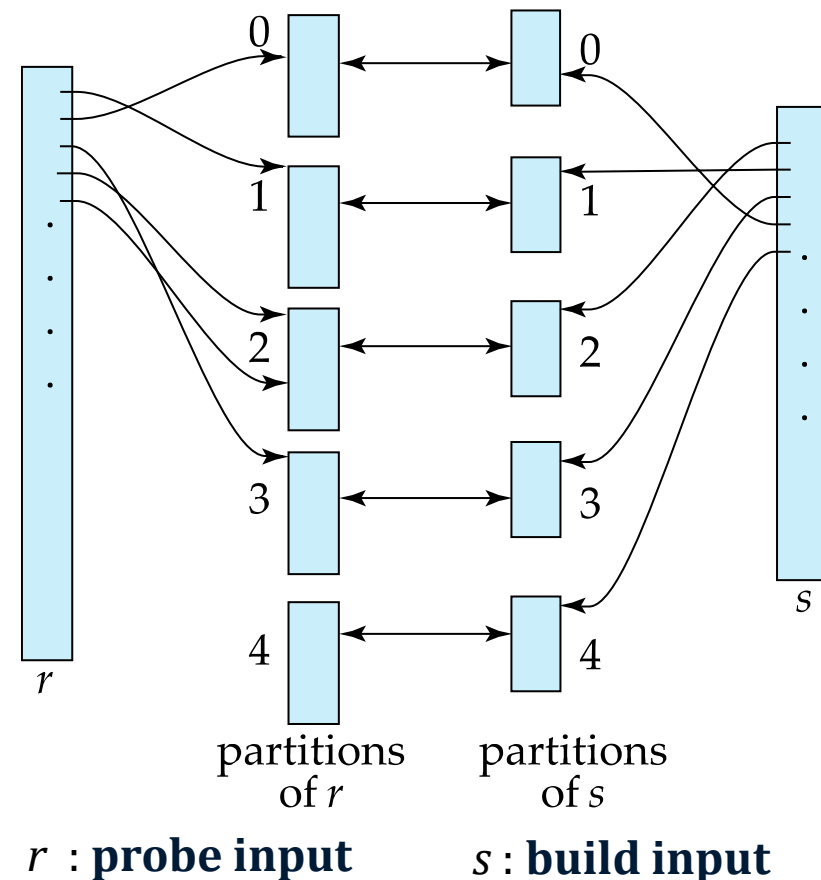
Step 1. **Partition** the relation s using a **hash function** h (into n_h partitions)

* note that n_h must be smaller than M

Step 2. **Partition** r similarly



Requirement: in step 1, we require that each partition of s must have at most $M-2$ blocks (i.e., can fit in memory)



r : **probe input**

s : **build input**

Hash-Join

Example: $n_h = 5$ partitions

Step 3. For each partition value i

(a) Create another hash function h'

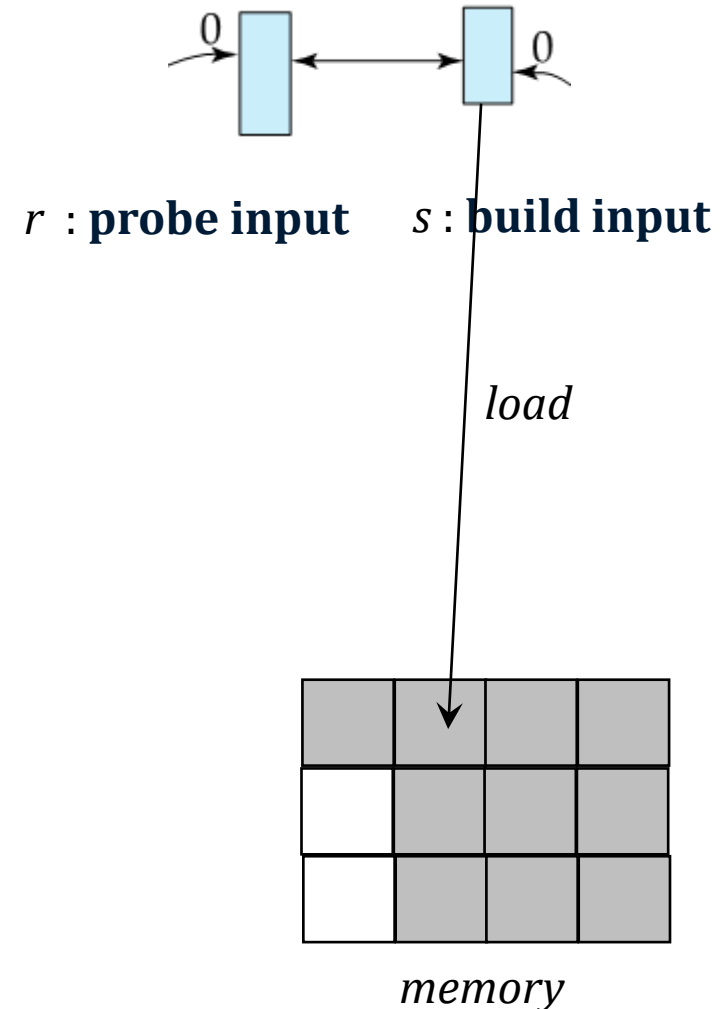
(b) Load the **entire** s_i into memory.

Build an in-memory hash index on it by the join attribute (by using h').

(c) Read the tuples in r_i .

For each tuple t_r , find each matching tuple t_s in s_i from in-memory hash index. Output the result.

Purpose of the in-memory hash index:
reduce the computation cost on
matching tuples



Hash-Join algorithm (Cont.)

- Suppose that number of partitions n_h is at most number of memory blocks M

- Cost of hash join:

$$3(b_r + b_s)$$

- Partitioning phase:

$$2(b_r + b_s)$$

- read and write each relation once

- Build and probe phase:

$$b_r + b_s$$

- Partially filled blocks:

$$2(2n_h)$$

small compared to other terms, can be ignored



- Recursive partitioning is not required if $M > n_h + 1$**

- E.g., recursive partitioning not needed for relations of 1GB with memory size of 2MB and block size of 4KB

Exercise: Cost of Hash-Join

customer ⋈ *depositor*

- ◆ Given that
 - ◆ memory size is 20 blocks
 - ◆ $b_{depositor} = 100$ and $b_{customer} = 400$.
- ◆ Use the smaller relation (*depositor*) as build input
- ◆ How large should a partition be?
 - ◆ To make each partition of *depositor* fit in memory (20 blocks), we can partition it into $\lceil 100/(20-2) \rceil = 6$ partitions
 - ◆ Since 6 < 20, this partitioning can be done in one pass
 - ◆ Similarly, we can divide *customer* into 6 partitions
- ◆ The total cost:
 - ◆ $3(100 + 400) = 1500$ blocks

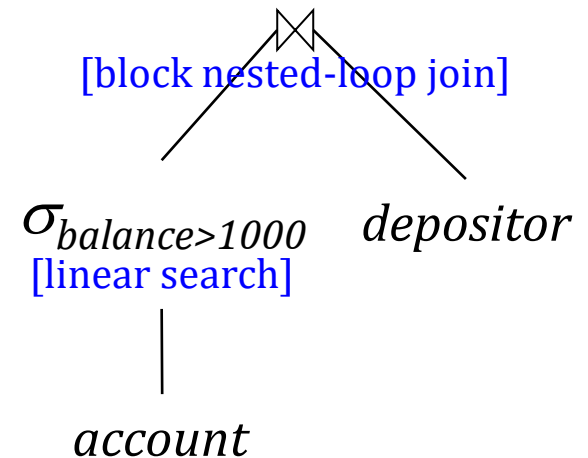
Outline

- ◆ Basic concepts for query processing
- ◆ How to process a selection?
- ◆ How to process a join?
- ◆ How to execute a plan?



How to execute a plan?

- ◆ **Naïve approach:** execute physical operators one-by-one
 - ◆ First execute “linear search” (for selection)
 - ◆ if the intermediate result exceeds the memory size, need to write it to the disk → additional I/O cost
 - ◆ Then execute “block nested-loop join” (for natural join)



- ◆ **Drawback**
 - ◆ High latency for the first result record
 - ◆ Additional I/O cost for writing intermediate result to the disk

How to execute a plan?

- ◆ The **iterator approach**: each physical operator implements the “Iterator” interface

- ◆ Open(): open the file/index, allocate buffer
- ◆ **GetNext()**: produce a record as output
- ◆ Close(): close the file/index, deallocate buffer

$\sigma_{balance > 1000}$
[linear search]

|
account

- ◆ *Example*: the iterator for linear search

B: 

Open ()


1. allocate a memory block *B*
2. open the file

GetNext ()

1. repeat
2. if *B* is empty, load the next block from disk to *B*
3. $t \leftarrow$ pop the next tuple from *B*
4. until *t* satisfies the predicate
5. return *t*

How to execute a plan?

- Advantages of the iterator approach
 - Low latency for the first result record
 - We can pipeline intermediate results to the next/parent operator, without additional I/O cost
- Iterators incur computational overhead
 - Acceptable in traditional RDBMS
- An issue in query optimization
 - How to allocate/share the memory to different physical operators?



Summary

- ◆ How to process a selection?
 - ◆ Linear scan, several types of index scan
- ◆ How to process a join?
 - ◆ Nested-loop, block nested-loop, block nested-loop, merge-join, hash-join
- ◆ Query plan execution

Sample solutions

- ◆ Please be reminded to submit Assignment 2 on time
- ◆ We will post the sample solutions of Quiz 2 and Assignment 2 at least one day before the exam

Sample types of questions

- ◆ Give an example of ... such that ...
- ◆ Draw ... such that ...
- ◆ Run ... show the steps ...
- ◆ Find ... show the steps ...
- ◆ Write an algorithm to ...
- ◆ Check ... explain ...

Wish you good luck! 😊