

Recall we want to maximize the following objective in RLHF

$$\mathbb{E}_{\hat{y} \sim p_{\theta}^{RL}(\hat{y}|x)} [RM_{\phi}(x, \hat{y}) - \beta \log \left(\frac{p_{\theta}^{RL}(\hat{y}|x)}{p^{PT}(\hat{y}|x)} \right)]$$

There is a closed form solution to this:

$$p^*(\hat{y}|x) = \frac{1}{Z(x)} p^{PT}(\hat{y}|x) \exp\left(\frac{1}{\beta} RM(x, \hat{y})\right)$$

- Rearrange this via a log transformation

$$RM(x, \hat{y}) = \beta (\log p^*(\hat{y}|x) - \log p^{PT}(\hat{y}|x)) + \beta \log Z(x) = \beta \log \frac{p^*(\hat{y}|x)}{p^{PT}(\hat{y}|x)} + \beta \log Z(x)$$

- This holds true for any arbitrary LMs, thus

$$RM_{\theta}(x, \hat{y}) = \beta \log \frac{p_{\theta}^{RL}(\hat{y}|x)}{p^{PT}(\hat{y}|x)} + \beta \log Z(x)$$

Putting It Together for DPO



✓ DPO: like finetune, direct use human data.
RLHF: train Reward Model, sample from reward model.

- Derived reward model: $RM_{\theta}(x, \hat{y}) = \beta \log \frac{p_{\theta}^{RL}(\hat{y}|x)}{p^{PT}(\hat{y}|x)} + \beta \log Z(x)$

- Final DPO loss via the Bradley-Terry model of human preferences:

Log Z term
cancels as
the loss only
measures
differences
in rewards

$$J_{DPO}(\theta) = -\mathbb{E}_{(x, y_w, y_l) \sim D} [\log \sigma(RM_{\theta}(x, y_w) - RM_{\theta}(x, y_l))]$$

$$= -\mathbb{E}_{(x, y_w, y_l) \sim D} \left[\log \sigma \left(\beta \log \frac{p_{\theta}^{RL}(y_w|x)}{p^{PT}(y_w|x)} - \beta \log \frac{p_{\theta}^{RL}(y_l|x)}{p^{PT}(y_l|x)} \right) \right]$$

DPO don't need to train

Reward Model like RLHF,

Reward for
winning sample

Reward for
losing sample

RLHF: $RM_{\psi}(x, \hat{y})$

about RM_{ψ} [Rafailov+ 2023] (PPO)


DPO: $\mathbb{E} \log \sigma(RM_{\theta}(x, y_w) - RM_{\theta}(x, y_l))$

about y_w and y_l (Adam)

DPO Performs Better



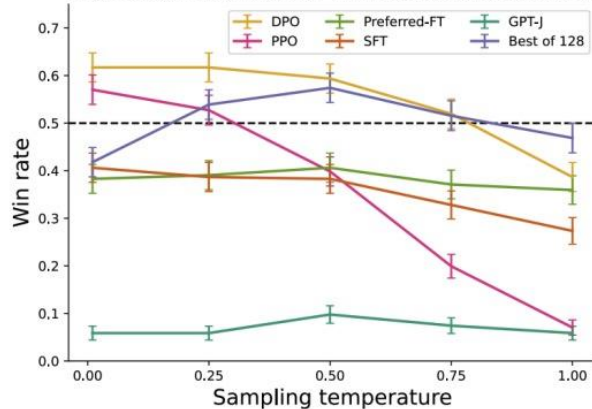
Reinforcement Learning from Human Feedback (RLHF)



Direct Preference Optimization (DPO)



TL;DR Summarization Win Rate vs Reference



- You can replace the complex RL part with a very simple weighted MLE objective
- Other variants (KTO, IPO) now emerging too
- TL;DR summarization win rates vs. human-written summaries (GPT-4 as a judge)

Assumption 3.2 f is a **convex function** and

$$\mathbb{E}_{\xi}[g(\theta, \xi)] = \nabla f(\theta),$$

$$\mathbb{E}_{\xi}[\|g(\theta, \xi)\|^2] \leq \textcircled{B^2} \quad \forall \theta.$$



where B is a given parameters. \hookrightarrow Variance Reduce Algorithms:
SVRG, SAGA

Theorem 3.2 Let $\{\theta^k\}$ be the sequence generated by SGD with step size $\alpha_k > 0$, under Assumption 3.2, for any $T > 0$,

$$\mathbb{E}[f(\bar{\theta}^T) - f^*] \leq \frac{\|\theta^0 - \theta^*\|^2 + \textcircled{B^2} \sum_{j=0}^T \alpha_j^2}{2 \sum_{j=0}^T \alpha_j}, \quad \downarrow \quad \triangle$$

where

$$\lambda_k = \sum_{j=0}^k \alpha_j, \quad \bar{\theta}^k = \lambda_k^{-1} \sum_{j=0}^k \alpha_j \theta^j.$$



Start from some $\theta^0 \in \mathbb{R}^s$, $v_0 = g(\theta^0, \xi_0)$, for $k \geq 0$:

$$\begin{aligned}\vartheta^k &= \theta^k - \beta_k v^k, \\ v^{k+1} &= \beta_k v^k + \alpha_k g(\vartheta^k, \xi_k), \\ \theta^{k+1} &= \theta^k - v^{k+1}.\end{aligned}$$

An advantage: prevent overshoot!



AdaGrad: Adaptive Learning Rates

Key idea: Rescale the learning rate of each coordinate by the historical progress.

2nd order moment: $E(X \cdot X^T) = \begin{pmatrix} E(x_1^2) & \dots & E(x_1 x_n) \\ \vdots & & \vdots \\ E(x_n x_1) & \dots & E(x_n^2) \end{pmatrix}$

Start from some $\theta^0 \in \mathbb{R}^s$, $n_g = 0$, for $k \geq 0$:

$$\begin{aligned} n_g &= n_g + g(\theta^k, \xi_k) \cdot * g(\theta^k, \xi_k), \\ \theta^{k+1} &= \theta^k - \alpha_k g(\theta^k, \xi_k) ./ (n_g + 10^{-8}). \end{aligned}$$

(n_g go to ∞ quickly)

Issue: The learning rate (step size) goes to zero quickly.

2nd order optimization:
 $\nabla^2 f(x) = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \vdots & & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{pmatrix}$ Hessian, Newton descent method.

Momentum + Adagrad

Key idea: Consider momentum and adaptive learning rate (second-order momentum) together.

Momentum + Adagrad

1st order moment + 2nd order moment

Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1)$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)

end while

return θ_t (Resulting parameters)

Empirical risk minimization: to find a function $f(\cdot)$ to minimize

$$\frac{1}{n} \sum_{i=1}^n L(f(X_i, \theta), Y_i)$$


over

$\mathcal{F} = \{f: f(x; \theta) \text{ is a neural network}$
 $\text{parameterized by } \theta \in \mathbb{R}^s \text{ outputs real values}\}.$

We expect

- Surrogate Loss function $L(\cdot, \cdot)$: continuous, smooth
- Neural network $f(\cdot; \theta)$: output continuous value
- The estimation easy to implement and explain

Feasible loss functions



0-1 loss:

Exponential loss (AdaBoost): $\phi(y \cdot f(x, \theta)) = \exp(-y \cdot f(x, \theta))$

Logistic loss : $\phi(y \cdot f(x, \theta)) = \log\{1 + \exp[-y \cdot f(x, \theta)]\}$

Hinge loss (SVM): $\phi(y \cdot f(x, \theta)) = \max\{1 - y \cdot f(x, \theta), 0\}$



A derivation of Xavier Initialization

Consider $y = w_1x_1 + w_2x_2 + \dots + w_nx_n$, x_i are i.i.d. with zero mean, w_i are i.i.d with zero mean.

Target: Compute $Var[y]$.

Lemma $Var[w_ix_i] = (E[w_i])^2Var[x_i] + (E[x_i])^2Var[w_i] + Var[w_i]Var[x_i]$.

Thus, $Var[w_ix_i] = Var[w_i]Var[x_i]$ and

$$Var[y] = Var[w_1x_1 + w_2x_2 + \dots + w_nx_n] = \sum_{i=1}^n Var[w_ix_i] = nVar[w_i]Var[x_i]$$

Thus,

$$Var[w_i] = 1/n$$


$$w_i \sim N(0, 1)/\sqrt{n}$$

He's Xavier Initialization for ReLu Activation Function



Key Motivation: Assume that only a half of the neurons are activated in each layer.


He's Xavier initialization: $N(0, 1/\sqrt{n_{in}/2})$.



Generative models




GAN: minimax the classification error loss.




probability model
VAE: maximize ELBO.

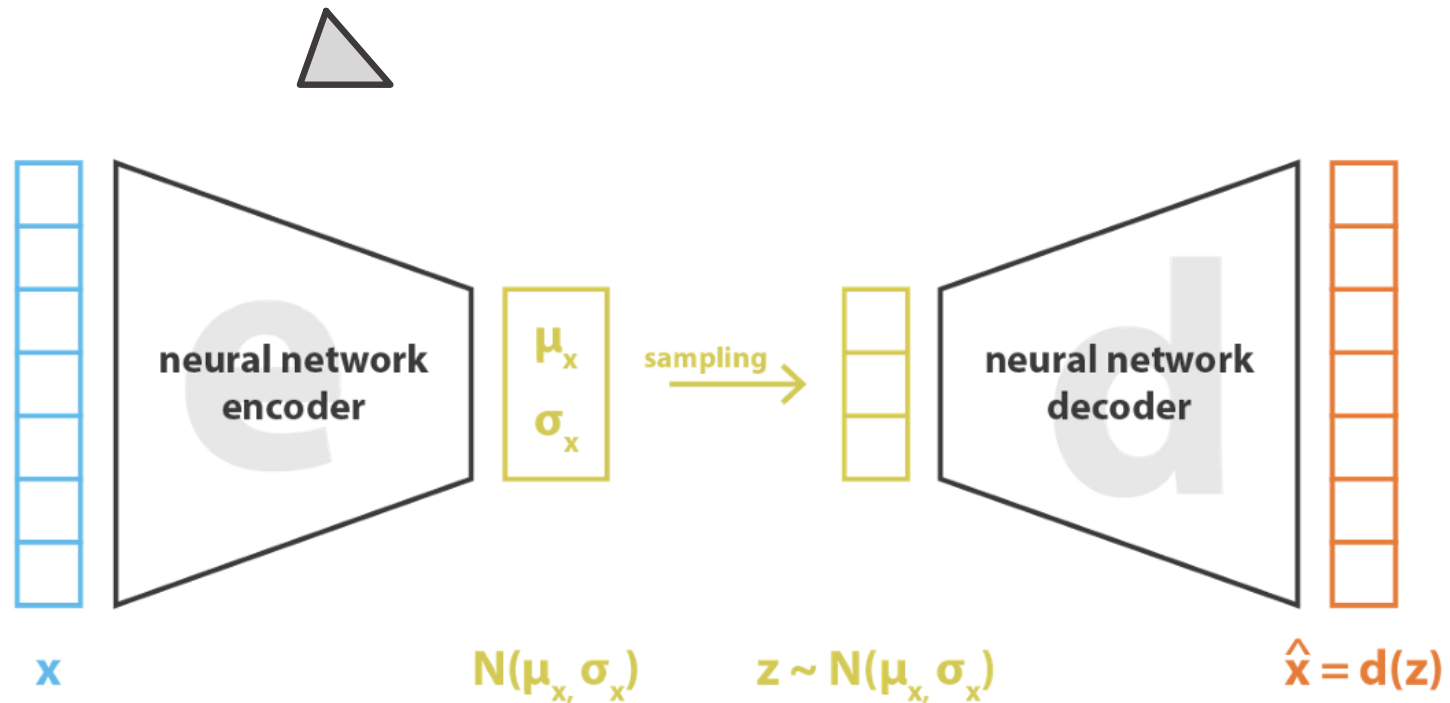
AE: compression model.



Flow-based generative models:
minimize the negative log-likelihood



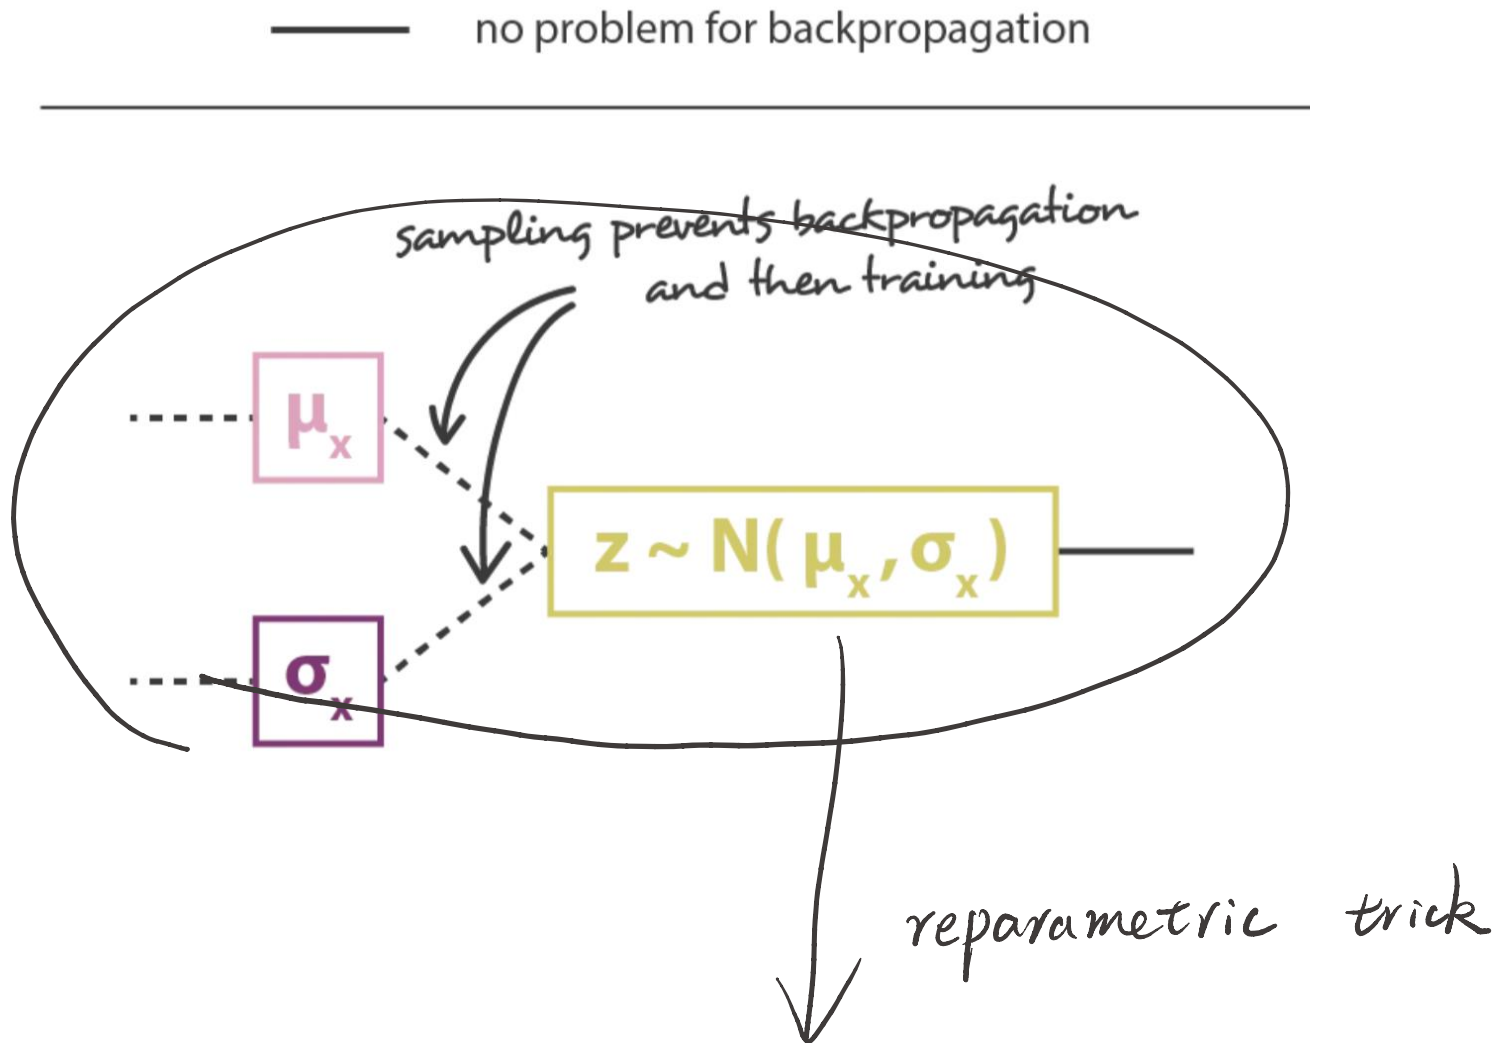
Source: <https://lilianweng.github.io/posts/2018-10-13-flow-models/>



$$\text{loss} = ||x - \hat{x}||^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)] = ||x - d(z)||^2 + \text{KL}[N(\mu_x, \sigma_x), N(0, I)]$$

The loss function is composed of
a **reconstruction term** and a **regularisation term**.

Sampling prevents backpropagation and the training

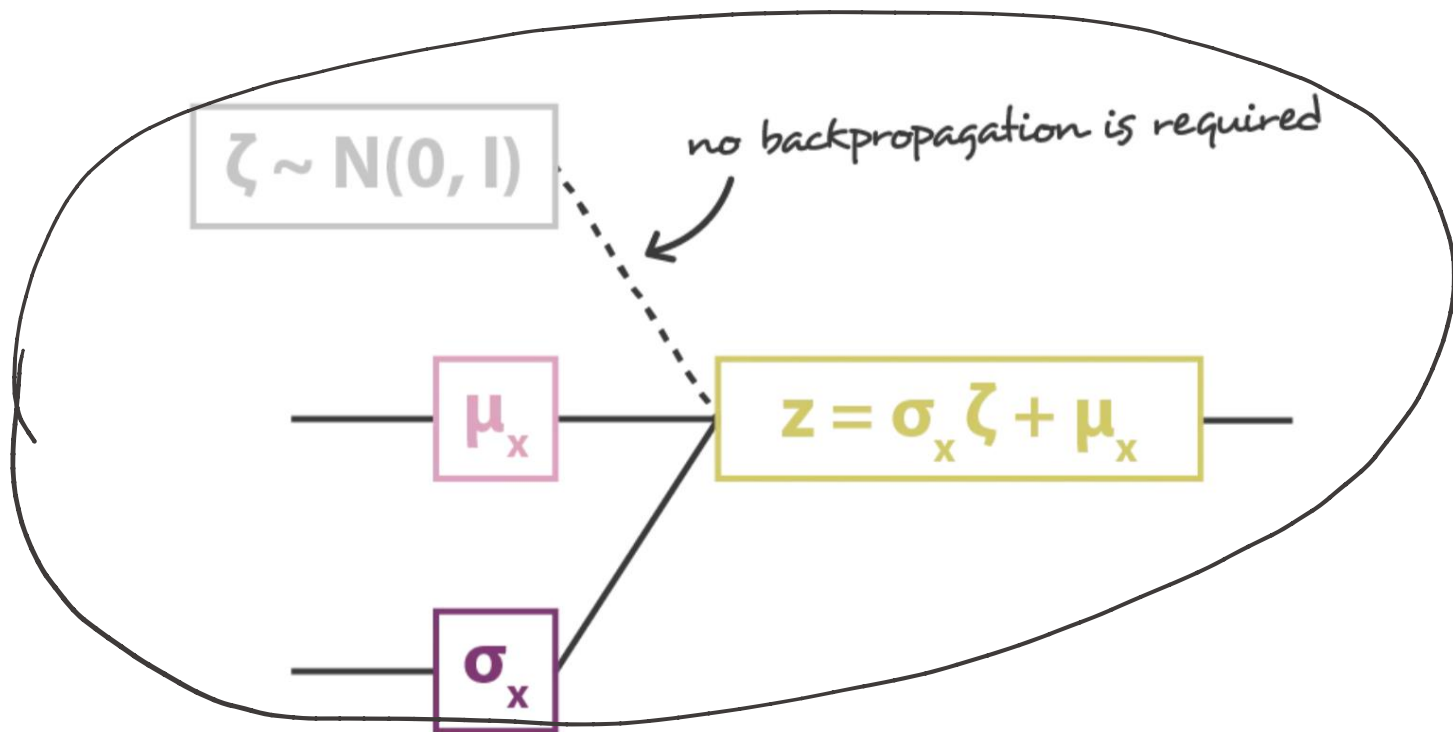


Reparameterization trick

$$z = \sigma_x \zeta + \mu_x$$

$$\zeta \sim N(0, I)$$

..... backpropagation is not possible due to sampling



Score-based models

diffusion model

- Train score-based models by minimizing the Fisher divergence

$$\mathbb{E}_{p(\mathbf{x})}[\|\nabla_{\mathbf{x}} \log p(\mathbf{x}) - \mathbf{s}_{\theta}(\mathbf{x})\|_2^2]$$

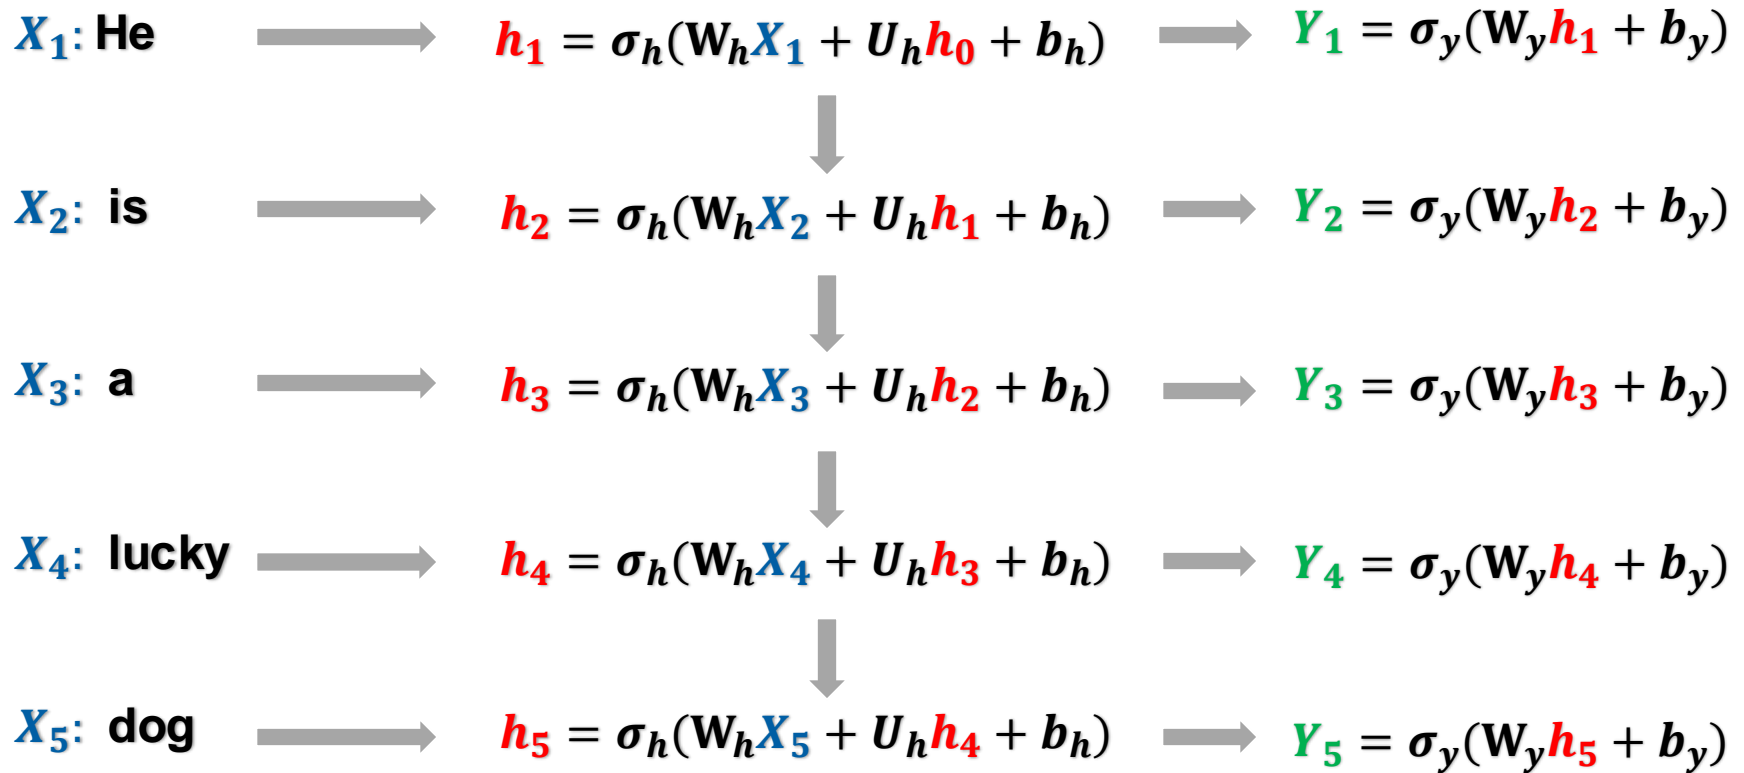
- Once trained a model $s_{\theta}(x) \approx \nabla_x \log p(x)$, we can use an iterative procedure called Langevin dynamics to draw samples from it.
- It initializes the chain from an arbitrary prior distribution $x_0 \sim \pi(x)$, and then iterates the following

$$\mathbf{x}_{i+1} \leftarrow \mathbf{x}_i + \epsilon \nabla_{\mathbf{x}} \log p(\mathbf{x}) + \sqrt{2\epsilon} \mathbf{z}_i, \quad i = 0, 1, \dots, K,$$


where \mathbf{z}_i follow standard Gaussian distribution.

- When $\epsilon \rightarrow 0$ and $K \rightarrow \infty$, x_K obtained from the procedure converges to a sample from $p(x)$ under some regularity conditions

Elman Network




Jordan Network



Long Short-Term Memory

The forward pass of an LSTM cell with a forget gate are



$$f_t = \sigma_g(W_f x_t + U_f h_{t-1} + b_f)$$

$$i_t = \sigma_g(W_i x_t + U_i h_{t-1} + b_i)$$

$$o_t = \sigma_g(W_o x_t + U_o h_{t-1} + b_o)$$

$$\tilde{c}_t = \sigma_c(W_c x_t + U_c h_{t-1} + b_c)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

$$h_t = o_t \odot \sigma_h(c_t)$$

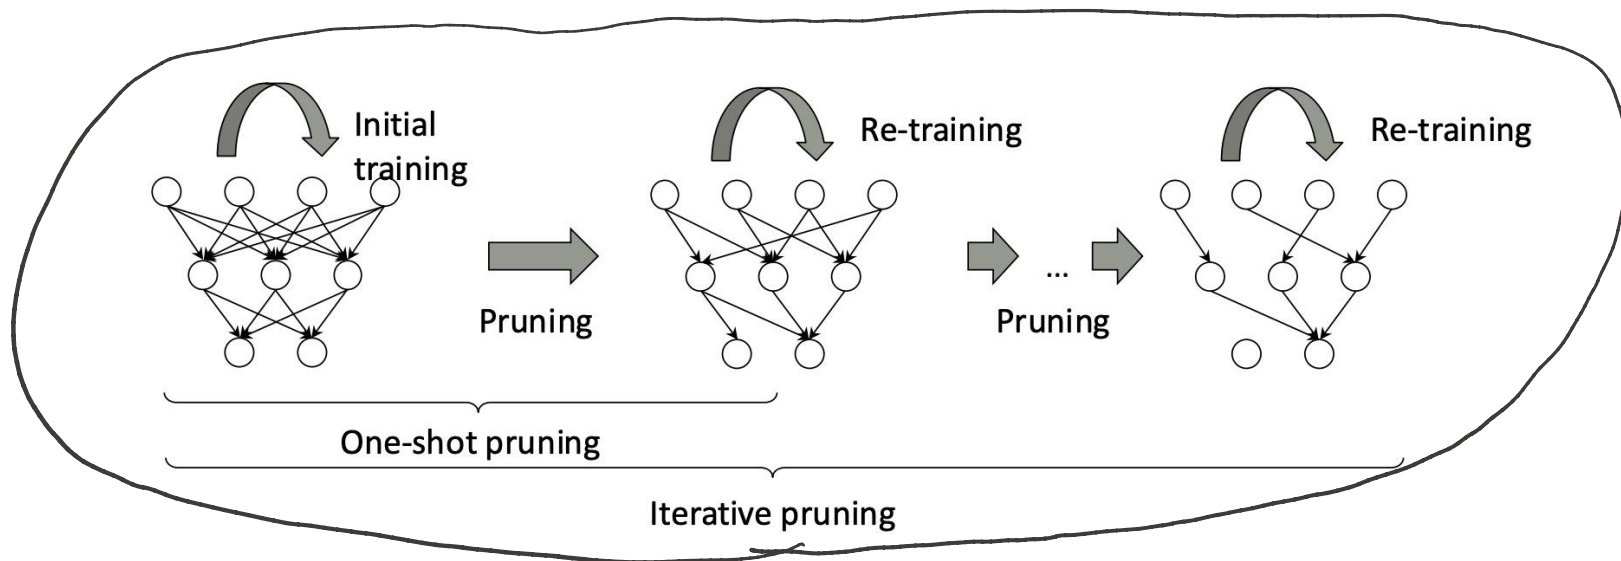
where the initial values are $c_0 = \mathbf{0}$ and $h_0 = \mathbf{0}$, and the operator \odot denotes the Hadamard product (element-wise product).

- $x_t \in \mathbb{R}^d$: input vector to the LSTM unit
- $f_t \in (0, 1)^h$: forget gate's activation vector
- $i_t \in (0, 1)^h$: input/update gate's activation vector
- $o_t \in (0, 1)^h$: output gate's activation vector
- $h_t \in (-1, 1)^h$: hidden state vector
- $\tilde{c}_t \in (-1, 1)^h$: cell input activation vector
- $c_t \in \mathbb{R}^h$: cell state vector

- **Continuous Skip-gram Model:**
predicts words within a certain range before and after the current word in the same sentence.
- **Continuous Bag-of-Words Model (CBOW):**
predicts the middle word based on surrounding context words. The context consists of a few words before and after the current (middle) word. This architecture is called a bag-of-words model as the order of words in the context is not important.

To train a neural network with a single hidden layer to perform a prediction task. But the goal is to learn the weights of the hidden layer—these weights are the “word vectors”.

- During pruning, a fraction of the lowest-magnitude weights are removed
- The non-pruned weights are re-trained
- Pruning for multiple iterations is more common ([Frankle & Carbin, 2019](#))



The Lottery Ticket Hypothesis

- Dense, randomly-initialized models **contain subnetworks** (“winning tickets”) that—when trained in isolation—**reach test accuracy comparable to the original network** in a similar number of iterations ([Frankle & Carbin, 2019](#))